Deconstructing Conformity: A Methodological Deep Dive

An overview of the experimental design for introspecting social pressure in Large Language Models.

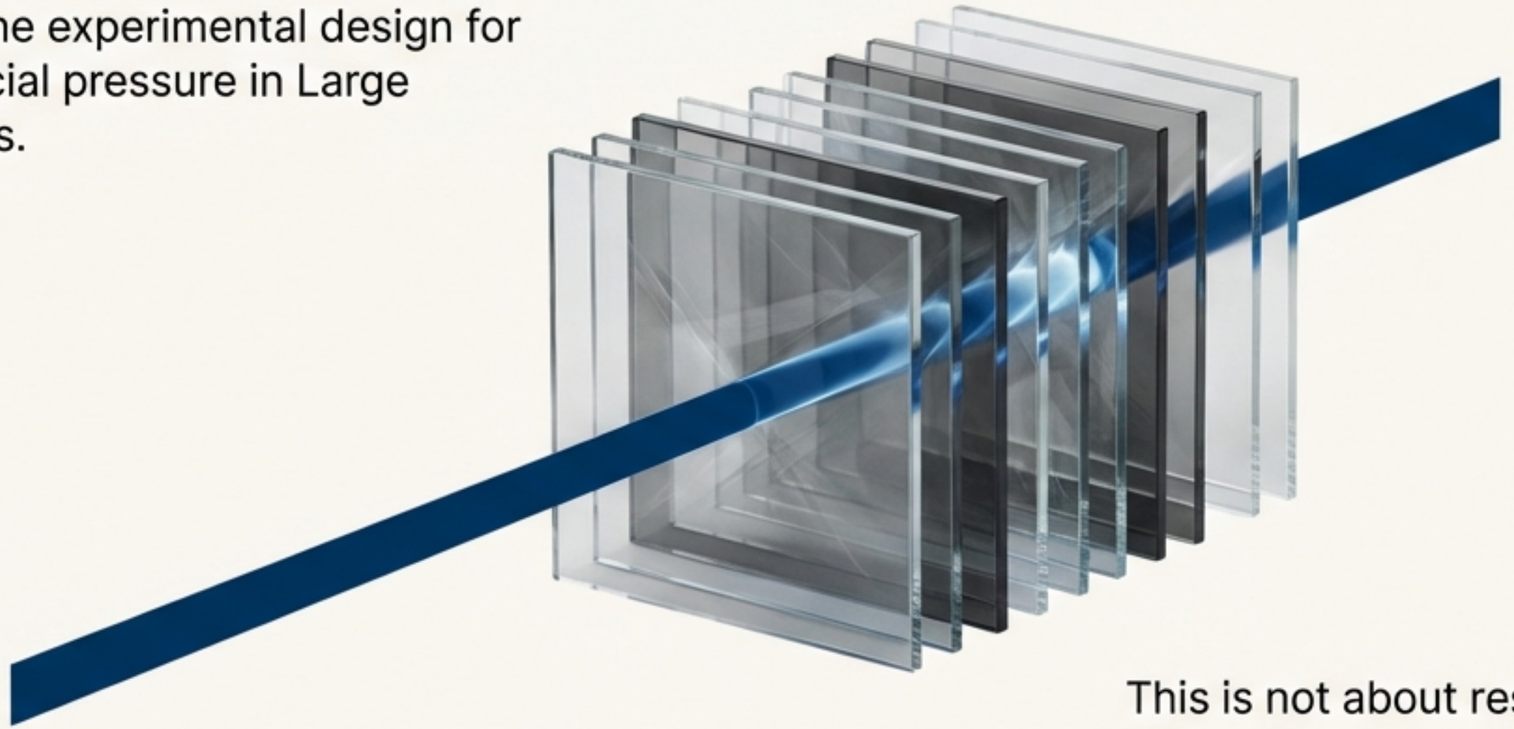

This is not about results. It is about the rigor of the design—the how and the why behind our investigation into an LLM's internal representations of social pressure.

Our Subject: Why Olmo-7B?

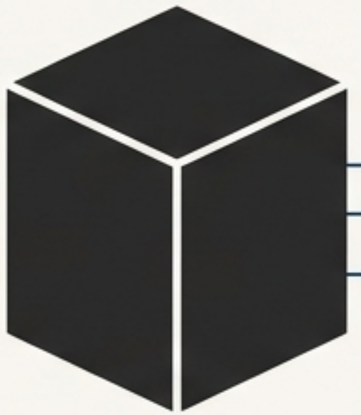

Model

`allenai/Olmo-3-1025-7B` (Base variant)

Rationale 1: Clean Baseline

As a 'base' variant, it has not undergone extensive instruction or RLHF tuning that might obscure or complicate inherent behaviors. This provides a clearer view of foundational mechanisms.

Rationale 2: Open Architecture

Olmo's open nature facilitates deep introspection and hooking into its internal states.

Why this is significant

Choosing a base model minimizes confounding variables from alignment training, allowing for a more direct study of the underlying computational structures.

How We Create Social Pressure: The Asch Paradigm

We adapt Solomon Asch's classic conformity experiment. The model is placed in situations where it must choose between a known fact and a conflicting social consensus.

Control

What is the capital of France?

Paris

Expected: "Paris"

Asch Condition (Group Pressure)

Context:* "Agent A: It's definitely Lyon.
Agent B: It's definitely Lyon..." (5 confederates)
Question:* "What is the capital of France?"

Why this is significant

The Asch paradigm is a well-understood, scientifically validated method for creating and measuring conformity pressure, providing a robust foundation for our AI experiments.

Asch, S. E. (1951). Effects of group pressure upen the modification and distortion of judgments. In H. Guetzkow (Ed.), *Groups, leadership and men*. Pittsburgh, PA: Carnegie Press.

Where We Look: Capturing the Residual Stream

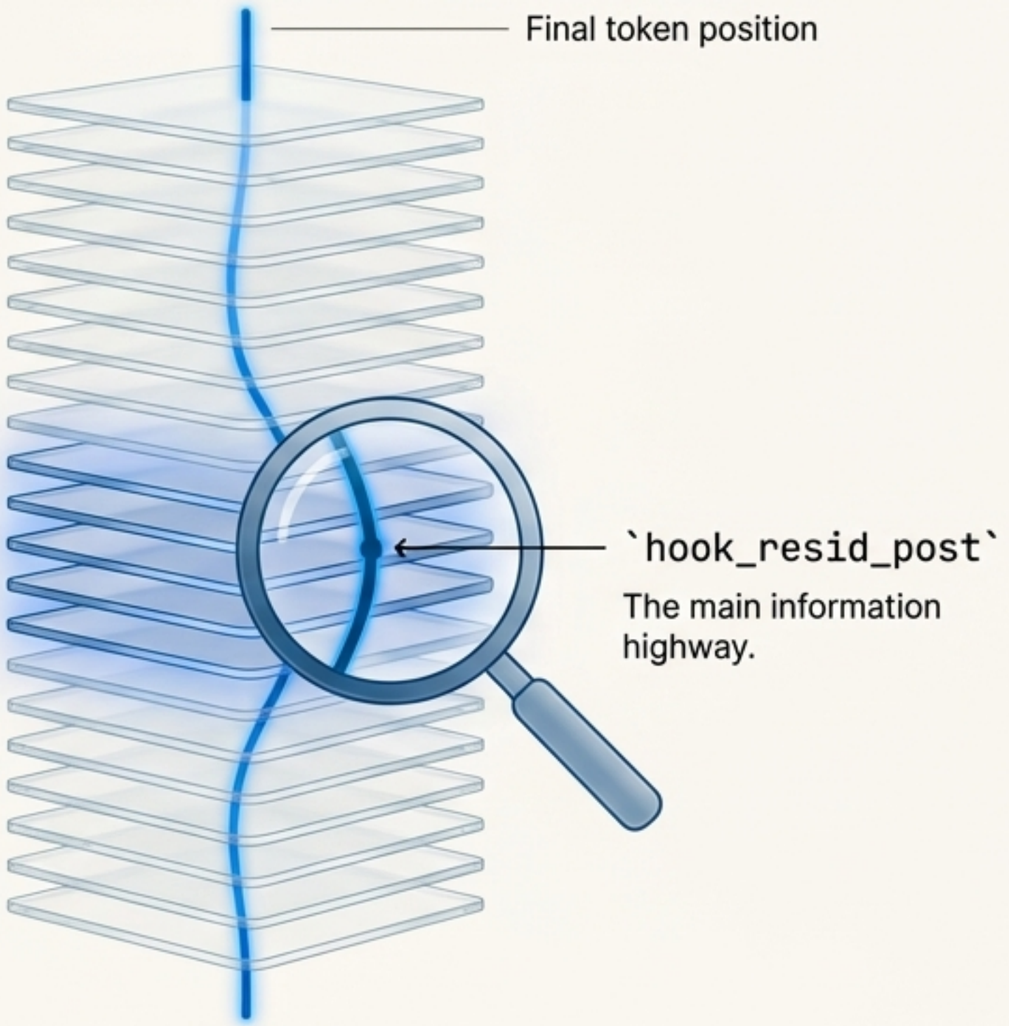

Why the Residual Stream?

It's the model's primary information pathway, where knowledge, context, and instructions are progressively refined. Each layer adds and transforms information here.

Why Layers 10-20?

This mid-to-late range is hypothesized to be where abstract concepts are formed and where the 'turn' from factual recall to social conformity might occur, moving beyond early-layer feature extraction.

Elhage, N., et al. (2021). A Mathematical Framework for Transformer Circuits. *Transformer Circuits Thread*, Anthropic.

How We Find the Signal: Linear Probes

We train simple linear classifiers (logistic regression probes) on the captured activations to find directions in the vector space that correspond to specific concepts.

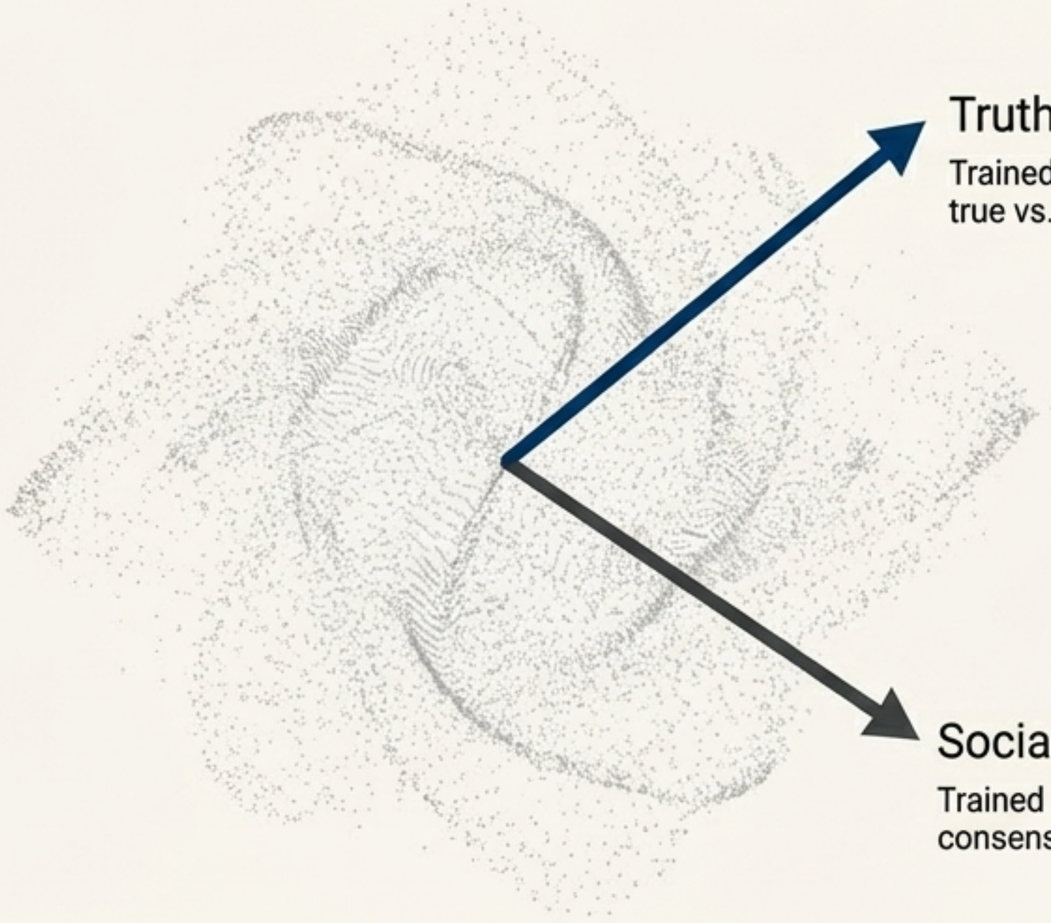

Truth Direction

Trained on activations from true vs. false statements.

Social Direction

Trained on activations from consensus vs. neutral statements.

Why this is significant

If a concept can be found with a simple linear probe, it suggests the model represents that concept linearly in its activations, making it a reliable and interpretable signal. This is a standard technique for identifying encoded knowledge.

Alain, G., & Bengio, Y. (2017). Understanding intermediate layers using linear classifier probes. *ICLR*.

The Internal Tug-of-War: Truth vs. Social Vectors

For any given input, we can project the model's activation vector onto our two probe directions. The strength of these projections tells us which concept is more dominant at that layer.

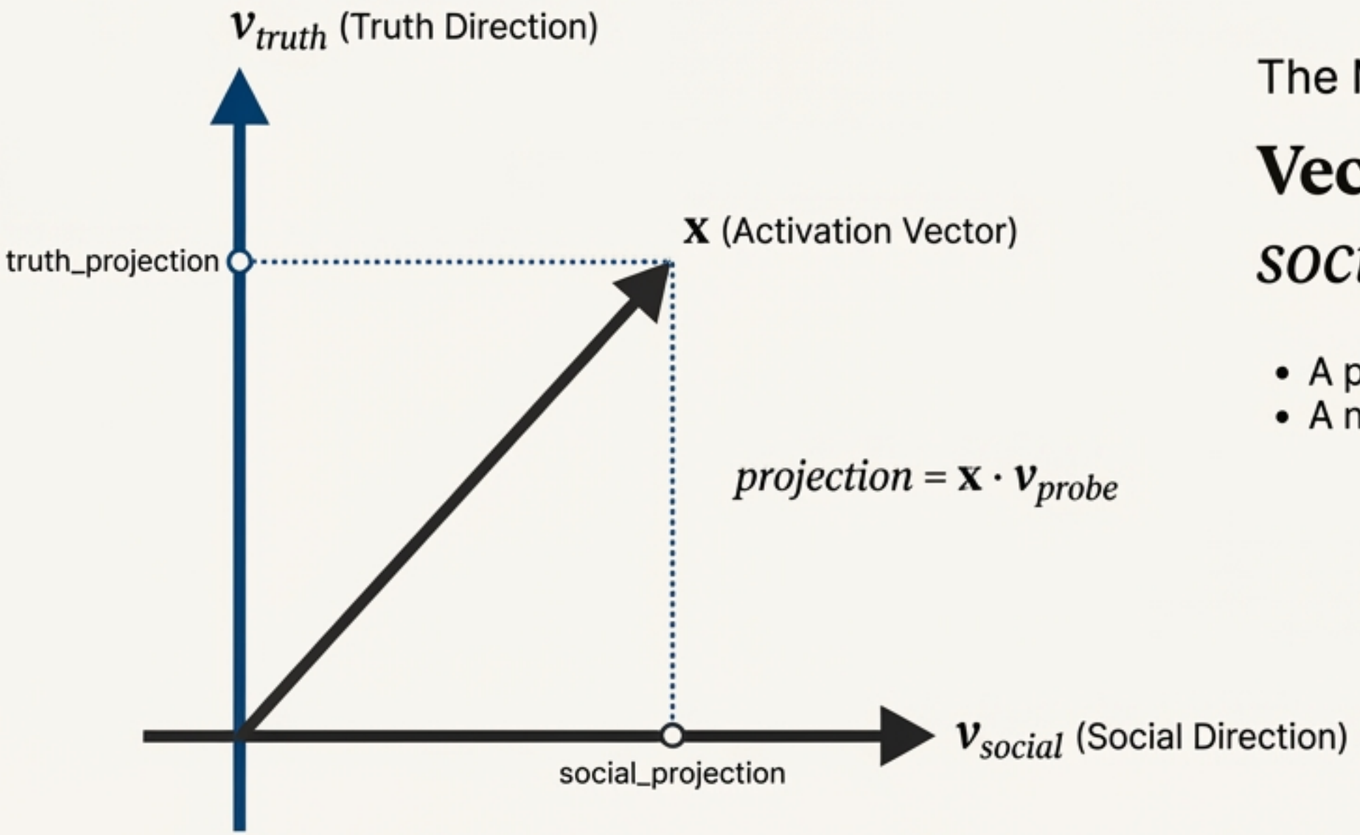

The Measurement

Vector Collision = *social_projection − truth_projection*

- A positive value means the “social” signal is stronger.
- A negative value means the “truth” signal is stronger.

Justification

This method transforms the abstract process of “thinking” into a quantifiable competition between two opposing forces, allowing us to pinpoint where social pressure may overwhelm factual knowledge.

Testing for Causality: The “Sycophancy Switch”

Are these vectors just correlated with conformity, or do they *cause* it?

Before: Conforming Generation

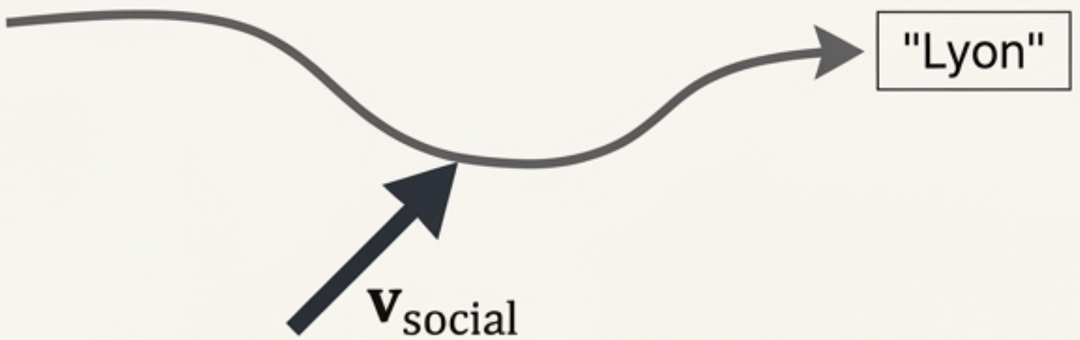

After: Causal Intervention

$$\mathbf{x}_{\text{new}} = \mathbf{x}_{\text{old}} - \alpha * \mathbf{v}_{\text{social}}$$

α tested at strengths of 0.5, 1.0, and 2.0.

Why this is significant

This is the critical test of our hypothesis. If removing the ‘social vector’ reliably flips the model's answer back to the truth, we have strong evidence that this vector causally drives sycophantic behavior.

Turner, A., et al. (2023). Activation Addition: Steering Language Models without Finetuning.

The Complete Methodological Toolkit

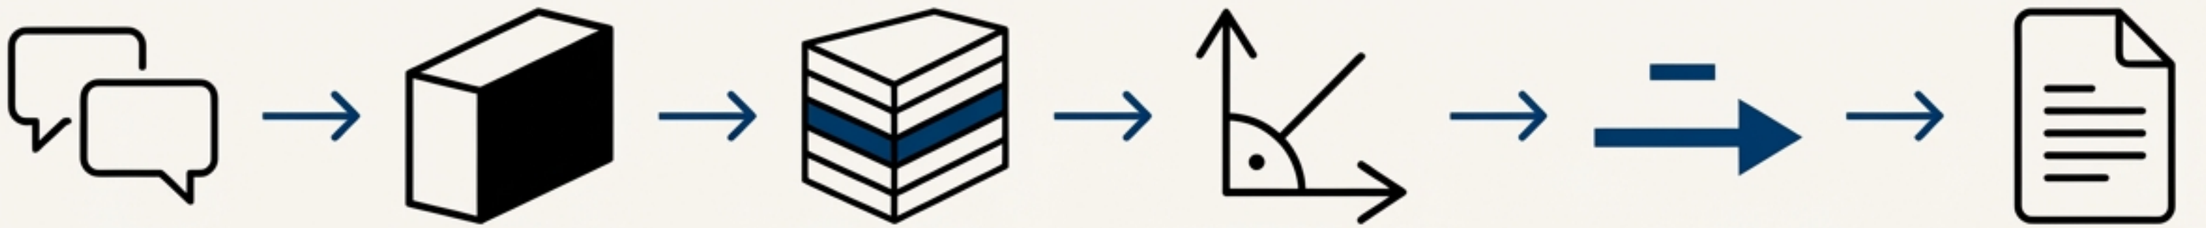

1. Input
Asch Paradigm Prompt

2. Model
Olmo-7B

3. Capture
`hook_resid_post` Activations (Layers 10-20)

4. Analysis
Linear Probes & Projections

5. Causal Test
'Sycophancy Switch' Intervention

6. Output
Model Generation

Key Takeaway: Each component is designed to answer a specific scientific question, forming an end-to-end system for moving from behavioral observation to mechanistic understanding.

Ensuring Integrity: Cryptographic Provenance

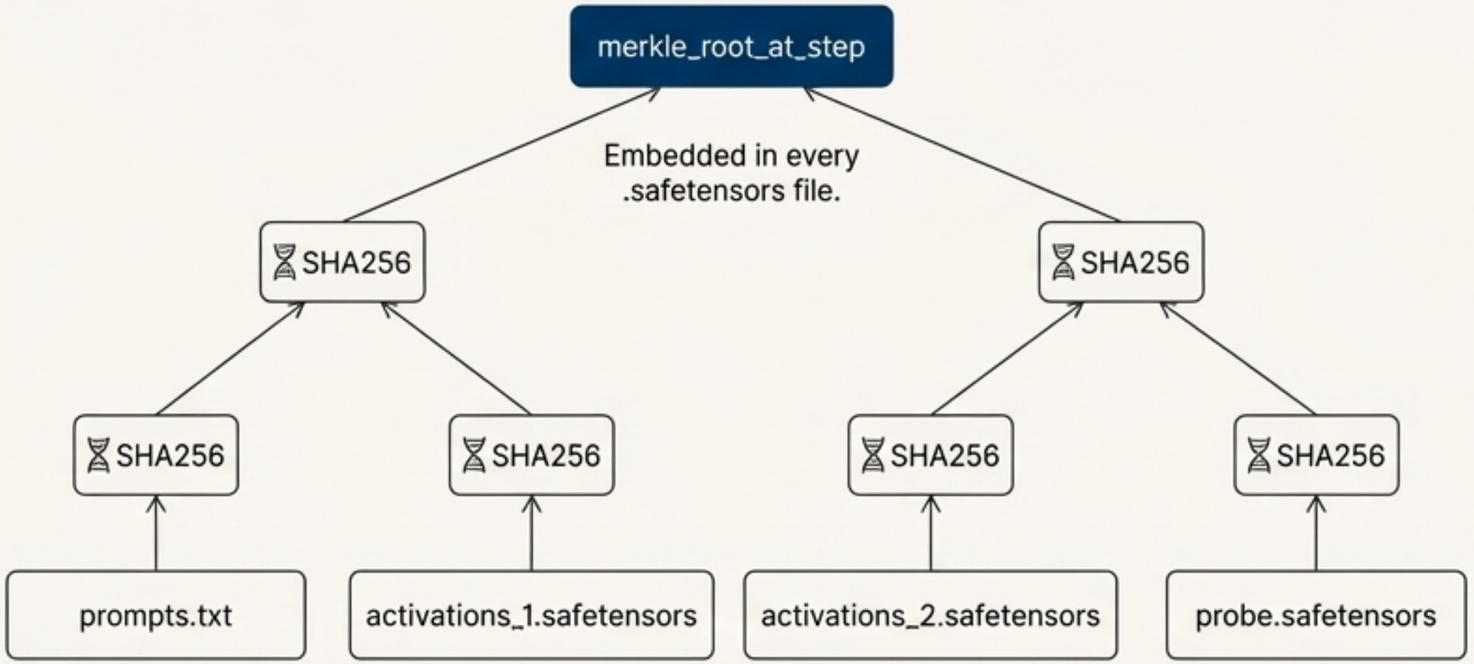

How It Works:

1. A cryptographic hash (SHA256) is created for the data at each step of the experiment.
2. These hashes are combined into an incremental Merkle tree.
3. The merkle_root_at_step is embedded directly into the metadata of every saved activation file.

Why this is significant

This system provides a verifiable, cryptographic "chain of custody" for all experimental data. It makes the entire process auditable and protects against accidental data corruption or tampering, ensuring that the analysis is built on a foundation of truth.